

Graph Algorithms Intro


- Problems on Graphs
- Cycle Checking
- Connected Components
- Hamiltonian Path and Circuit
- Euler Path and Circuit

❖ Cycle Checking

A graph has a cycle if

- it has a path of length > 2
- with start vertex src = end vertex $dest$ 起点 = 终点
- and without using any edge more than once


This graph has 3 distinct cycles: 0-1-2-0, 2-3-0-2, 0-1-2-3-0



("distinct" means the set of vertices on the path, not the order)

❖ ... Cycle Checking

Consider this graph:



This graph has many cycles e.g. 0-4-3-0, 2-4-5-2, 0-1-2-5-4-6-3-0, ...

❖ ... Cycle Checking

First attempt at checking for a cycle

```


hasCycle(G):
    Input  graph G
    Output true if G has a cycle, false otherwise

    choose any vertex  $v \in G$ 
    return dfsCycleCheck(G,v)

dfsCycleCheck(G,v):
    mark v as visited
    for each (v,w)  $\in$  edges(G) do
        if w has been visited then // found cycle
            return true
        else if dfsCycleCheck(G,w) then
            return true
    end for
    return false // no cycle at v
  
```

Handwritten notes: $v \rightarrow w$ (above the loop), $v \rightarrow w$ (above the recursive call), and *return true* (next to the recursive call's return statement).

COMP2521 20T2 ◊ Graph Algorithms [4/26]




❖ ... Cycle Checking


The above algorithm has two bugs...

- only one connected component is checked
- the loop **for each $(v,w) \in \text{edges}(G)$ do** should exclude the neighbour of v from which you just came, so as to prevent a single edge $w-v$ being classified as a cycle.

If we start from vertex 5 in the following graph, we don't find the cycle:




Connected Component #1



Connected Component #2

Bug


 不能用一个 path

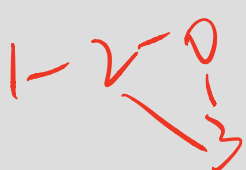
❖ ... Cycle Checking

Version of cycle checking (in C) for one connected component:

```

bool dfsCycleCheck(Graph g, Vertex v, Vertex u) {
    visited[v] = true;
    for (Vertex w = 0; w < g->nV; w++) {
        if (adjacent(g, v, w)) {
            if (!visited[w]) { 没走过w
                if (dfsCycleCheck(g, w, u) u to w)
                    return true;
            }
            else if (w != u) 走过w,
                return true;
        }
    }
    return false;
}

```



v=1 u=1 w=2
1
v=2
w=0 true
w=1 走过, w=u false
w=3 true

❖ ... Cycle Checking

Wrapper to ensure that all connected components are checked:

```
Vertex *visited;  
  
bool hasCycle(Graph g, Vertex s) {  
    bool result = false;  
    visited = calloc(g->nV, sizeof(int));  
    for (int v = 0; v < g->nV; v++) {  
        for (int i = 0; i < g->nV; i++)  
            visited[i] = -1;  
        if dfsCycleCheck(g, v, v) {  
            result = true;  
            break;  
        }  
    }  
    free(visited);  
    return result;  
}
```


❖ Connected Components

Consider these problems:

- how many connected subgraphs are there?
- are two vertices in the same connected subgraph?

Both of the above can be solved if we can

- build **componentOf[]** array, one element for each vertex v
- indicating which connected component v is in



`nComponents(g) = 3`

`componentOf[1] = 0`

`componentOf[5] = 2`

`sameComponent(3,4) = true`

`sameComponent(3,5) = false`

`sameComponent(0,6) = false`


❖ ... Connected Components

Algorithm to assign vertices to connected components:

```
components(G):  
  Input  graph G  
  Output componentOf[] filled for all v  
  
  for all vertices v ∈ G do  
    | componentOf[v] = -1  
  end for  
  compID = 0 // component ID  
  for all vertices v ∈ G do  
    | if componentOf[v] = -1 then  
      | dfsComponent(G, v, compID)  
      | compID = compID + 1  
    | end if  
  end for
```

DFS scan of one connected component

```
dfsComponent(G, v, id):  
  componentOf[v] = id  
  for each (v, w) ∈ edges(G) do  
    | if componentOf[w] = -1 then  
      | dfsComponent(G, w, id)  
    | end if  
  end for
```



❖ ... Connected Components

Consider an application where connectivity is critical

- we frequently ask questions of the kind above
- but we cannot afford to run **components()** each time

Add a new fields to the **GraphRep** structure:

```
typedef struct GraphRep *Graph;

struct GraphRep {
    ...
    int nC; // # connected components
    int *cc; // which component each vertex is contained in
    ...    // i.e. array [0..nV-1] of 0..nC-1
}
```

❖ ... Connected Components

With this structure, the above tasks become trivial:

```
// How many connected subgraphs are there?
int nConnected(Graph g) {
    return g->nC;
}

// Are two vertices in the same connected subgraph?
bool inSameComponent(Graph g, Vertex v, Vertex w) {
    return (g->cc[v] == g->cc[w]);
}
```

But ... introduces overheads ... maintaining **cc[]**, **nC**

❖ ... Connected Components

Consider maintenance of such a graph representation:

- initially, $nC = nV$ (because no edges)
- adding an edge may reduce nC
(adding edge between v and w in different components)
- removing an edge may increase nC
(removing edge between v and w in same component)
- $cc[]$ can simplify path checking
(ensure v, w are in same component before starting search)

Additional cost amortised by lower cost for $nConnected()$ and $inSameComponent()$

Is it simpler to run $components()$ after each edge change?

❖ Hamiltonian Path and Circuit

Hamiltonian path problem:

- find a simple path connecting two vertices v, w in graph G
- such that the path includes each vertex exactly once

$v \rightarrow v$
If $v = w$, then we have a Hamiltonian circuit

Simple to state, but difficult to solve (NP -complete)


Many real-world applications require you to visit all vertices of a graph:

- Travelling salesman
- Bus routes go to every stop.
- ...

Named after Irish mathematician/physicist/astronomer Sir William Hamilton (1805-1865)

❖ ... Hamiltonian Path and Circuit

Graph and some possible Hamiltonian paths:



要过所有的 V 一遍

❖ ... Hamiltonian Path and Circuit

Approach:

- generate all possible simple paths (using e.g. DFS)
- keep a counter of vertices visited in current path
- stop when find a path containing V vertices

Can be expressed via a recursive DFS algorithm

- similar to simple path finding approach, except
 - keeps track of path length; succeeds if length = V
 - resets "visited" marker after unsuccessful path

❖ ... Hamiltonian Path and Circuit

Algorithm for finding Hamiltonian path:

```
visited[] // array [0..nV-1] to keep track of visited vertices

hasHamiltonianPath(G,src,dest):
    Input  graph G, plus src/dest vertices
    Output true if Hamiltonian path src...dest,
           false otherwise

    for all vertices v ∈ G do
        visited[v]=false
    end for
    return hamiltonR(G,src,dest,#vertices(G)-1)
```

0 → 2 3

```
hamiltonR(G,v,dest,d):
    Input  G      graph
           v      current vertex considered
           dest   destination vertex
           d      distance "remaining" until path found


    if v=dest then
        if d=0 then return true else return false
    else
        visited[v]=true
        for each (v,w) ∈ edges(G) where not visited[w] do
            if hamiltonR(G,w,dest,d-1) then
                return true
            end if
        end for
    end if
    visited[v]=false // reset visited mark
    return false
```

回到起点并 cover all vertex

❖ ... Hamiltonian Path and Circuit

Analysis: worst case requires $(V-1)!$ paths to be examined

Consider a graph with isolated vertex and the rest fully-connected



Checking **hasHamiltonianPath($g, x, 0$)** for any x

- requires us to consider every possible path
- e.g 1-2-3-4, 1-2-4-3, 1-3-2-4, 1-3-4-2, 1-4-2-3, ...
- starting from any x , there are $3!$ paths $\Rightarrow 4!$ total paths
- there is no path of length 5 in these $(V-1)!$ possibilities

There is no known simpler algorithm for this task \Rightarrow NP-hard.


Note, however, that the above case could be solved in constant time if we had a fast check for 0 and x being in the same connected component

❖ Euler Path and Circuit


Euler path problem:

- find a path connecting two vertices v, w in graph G
- such that the path includes each edge exactly once
(note: the path does not have to be simple \Rightarrow can visit vertices more than once)

If $v = w$, then we have an Euler circuit



Euler Path: 4-3-1-5-2-1



Euler Circuit: 1-2-5-4-3-1

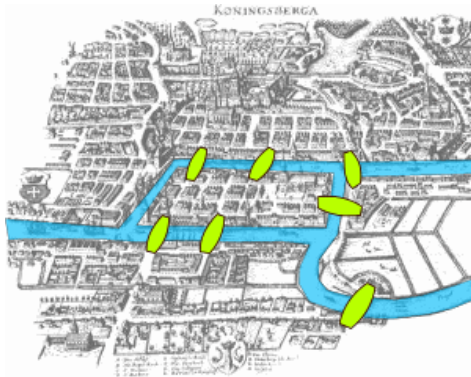
Many real-world applications require you to visit all edges of a graph:

- Postman *go every street*
- Garbage pickup
- ...

❖ ... Euler Path and Circuit

Problem named after Swiss mathematician, physicist, astronomer, logician and engineer Leonhard Euler (1707 - 1783)

Based on a circuitous route via bridges in Königsberg



❖ ... Euler Path and Circuit

One possible "brute-force" approach:

- check for each path if it's an Euler path
- would result in factorial time performance

Can develop a better algorithm by exploiting:


Theorem. A graph has an Euler circuit if and only if it is connected and all vertices have even degree




Theorem. A graph has a non-circuitous Euler path if and only if it is connected and exactly two vertices have odd degree

❖ ... Euler Path and Circuit

Graphs with an Euler path are often called Eulerian Graphs



Has neither Eulerian path or circuit



Has no Eulerian circuit, but does have path

5 - 4 - 3 - 0 - 1 - 2 - 3

❖ ... Euler Path and Circuit

Assume the existence of **degree(g, v)**

Algorithm to check whether a graph has an Euler path:

```
hasEulerPath(G,src,dest):
```

```
  Input graph G, vertices src,dest
```

```
  Output true if G has Euler path from src to dest  
         false otherwise
```

```
  if src≠dest then
```

```
    if degree(G,src) is even or degree(G,dest) is even then  
      return false
```

```
    end if
```

```
  end if
```


```
  for all vertices v ∈ G do
```

```
    if v≠src ∧ v≠dest ∧ degree(G,v) is odd then  
      return false and
```

```
    end if
```

```
  end for
```

```
  return true
```



❖ ... Euler Path and Circuit

Analysis of **hasEulerPath** algorithm:

- assume that connectivity is already checked
- assume that **degree()** is available via $O(1)$ lookup
- single loop over all vertices $\Rightarrow O(V)$

If degree requires iteration over vertices

- cost to compute degree of a single vertex is $O(V)$
- overall cost is $O(V^2)$

\Rightarrow problem tractable, even for large graphs (unlike Hamiltonian path problem)

For the keen, a linear-time (in the number of edges, E) algorithm to compute an Euler path is described in [Sedgewick] Ch.17.7.

Directed/Weighted Graphs

- Generalising Graphs
- Directed Graphs (Digraphs)
- Digraph Representation
- Weighted Graphs
- Weighted Graph Representation
- Weighted Graph Implementation

❖ Generalising Graphs

Discussion so far has considered graphs as

- V = set of vertices, E = set of edges

Real-world applications require more "precision"

- some edges are directional (e.g. one-way streets)
- some edges have a cost (e.g. distance, traffic)


We need to consider directed graphs and weighted graphs

❖ Directed Graphs (Digraphs)

Directed graphs are ...


- graphs with V vertices, E edges (v,w)
- edge (v,w) has source v and destination w
- unlike undirected graphs, $v \rightarrow w \neq w \rightarrow v$

Example digraph:



❖ ... Directed Graphs (Digraphs)

Some properties of ...



- edges 1-2-3 form a cycle, edges 1-3-4 do *not* form a cycle
- vertex 5 has a self-referencing edge (5,5)
- vertices 0 and 1 reference each other, i.e. (0,1) and (1,0)
- there are no paths from 5 to any other nodes
- paths from 0→5: 0-1-2-3-4-5, 0-1-4-5, 0-1-2-3-1-4-5

❖ ... Directed Graphs (Digraphs)

Terminology for digraphs ...

Directed path: sequence of $n \geq 2$ vertices $v_1 \rightarrow v_2 \rightarrow \dots \rightarrow v_n$

- where $(v_i, v_{i+1}) \in \text{edges}(G)$ for all v_i, v_{i+1} in sequence

If $v_1 = v_n$, we have a **directed cycle**

Degree of vertex: number of incident edges

- **outdegree:** $\deg(v)$ = number of edges of the form $(v, _)$ v →
- **indegree:** $\deg^{-1}(v)$ = number of edges of the form $(_, v)$ → v

❖ ... Directed Graphs (Digraphs)

More terminology for digraphs ...

Reachability:

- w is reachable from v if \exists directed path v, \dots, w

Strong connectivity:

- every vertex is reachable from every other vertex

Directed acyclic graph (DAG):

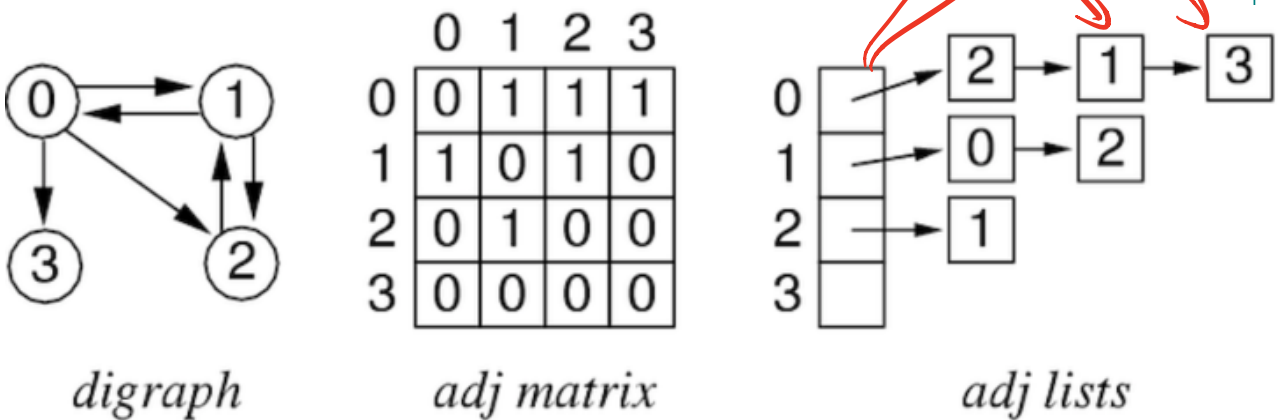
- contains no directed cycles

❖ Digraph Representation

Similar set of choices as for undirectional graphs:


- array of edges (directed)
- vertex-indexed adjacency matrix (non-symmetric)
- vertex-indexed adjacency lists

Vertices identified by $0..V-1$



❖ ... Digraph Representation

Example digraph and adjacency matrix representation:



	a	b	c	d	e	f	g
a	1	0	0	1	0	0	0
b	1	0	1	0	0	0	0
c	0	1	0	1	0	0	1
d	0	1	0	0	0	1	1
e	0	0	0	1	0	0	0
f	1	0	0	1	0	0	1
g	0	0	0	0	0	0	1


Undirectional \Rightarrow symmetric matrix

Directional \Rightarrow non-symmetric matrix

Maximum #edges in a digraph with V vertices: V^2


❖ ... Digraph Representation

Costs of representations: (where degree $\deg(v)$ = #edges leaving v)

	array of edges	adjacency matrix	adjacency list
space usage	E	V^2 	$V+E$
insert edge	E	1	1
exists edge (v,w) ?	E	1	$\deg(v)$
get edges leaving v	E	V	$\deg(v)$



Overall, adjacency list representation is best

- real graphs tend to be sparse 
(large number of vertices, small average degree $\deg(v)$)
- algorithms frequently iterate over edges from v

❖ Weighted Graphs

Graphs so far have considered

- edge = an association between two vertices/nodes
- may be a precedence in the association (directed)

Some applications require us to consider

- a cost or weight of an association
- modelled by assigning values to edges (e.g. positive reals)


❖ ... Weighted Graphs

Weighted graphs are ...


- graphs with V vertices, E edges (s,t)
- each edge (s,t,w) connects vertices s and t and has weight w

Weights can be used in both directed and undirected graphs.

Example weighted graphs:




Weighted Graph



Directed Weighted Graph

❖ ... Weighted Graphs

Example: major airline flight routes in Australia



Representation: edge = direct flight; weight = approx flying time (hours)

❖ ... Weighted Graphs


Weights lead to minimisation-type questions, e.g.

1. Cheapest way to connect all vertices?

- a.k.a. minimum spanning tree problem
- assumes: edges are weighted and undirected

2. Cheapest way to get from *A* to *B*?

- a.k.a. shortest path problem
- assumes: edge weights positive, directed or undirected



❖ Weighted Graph Representation

Weights can easily be added to:

- adjacency matrix representation ($0/1 \rightarrow \text{int or float}$)
- adjacency lists representation (add int/float to list node)

The edge list representation changes to list of (s, t, w) triples


All representations can also work with directed edges

Weight values are determined by domain being modelled

- in some contexts weight could be zero or negative

❖ ... Weighted Graph Representation

Adjacency matrix representation with weights:



Weighted Digraph


	0	1	2	3	4
0	*	0.2	0.4	*	*
1	*	0.3	0.6	*	*
2	*	0.5	*	0.1	*
3	0.5	*	*	*	0.9
4	*	*	0.1	*	*

Adjacency Matrix


Note: need distinguished value to indicate "no edge".

❖ ... Weighted Graph Representation

Adjacency lists representation with weights:



Weighted Digraph




Adjacency Lists


Note: if undirected, each edge appears twice with same weight

❖ ... Weighted Graph Representation

Edge array / edge list representation with weights:



Weighted Digraph



Edge List

Note: not very efficient for use in processing algorithms, but does give a possible representation for min spanning trees or shortest paths

❖ Weighted Graph Implementation

Changes to previous graph data structures to include weights:

WGraph.h

```
// edges are pairs of vertices (end-points) plus weight
typedef struct Edge {
    Vertex v;
    Vertex w;
    int weight;
} Edge;

// returns weight, or 0 if vertices not adjacent
int adjacent(Graph, Vertex, Vertex); 有边 edge
```

Note: here, we assume all weights are positive, but not required

❖ ... Weighted Graph Implementation

WGraph.c (assuming adjacency matrix representation)

```
typedef struct GraphRep {  
    int **edges; // adjacency matrix storing weights  
                // 0 if nodes not adjacent  
    int nV; // #vertices  
    int nE; // #edges  
} GraphRep;  
  
bool adjacent(Graph g, Vertex v, Vertex w) {  
    assert(valid_graph, valid_vertices)  
    return (g->edges[v][w] != 0);  
}
```

❖ ... Weighted Graph Implementation

More **WGraph.c**

```
void insertEdge(Graph g, Edge e) {
    assert(valid_graph, valid_edge)
    // edge e not already in graph
    if (g->edges[e.v][e.w] == 0) {g->nE++;}
    // may change weight of existing edge
    g->edges[e.v][e.w] = e.weight;
    g->edges[e.w][e.v] = e.weight;
}

void removeEdge(Graph g, Edge e) {
    assert(valid_graph, valid_edge)
    // edge e not in graph
    if (g->edges[e.v][e.w] == 0) {return;}
    g->edges[e.v][e.w] = 0;
    g->edges[e.w][e.v] = 0;
    g->nE--;
}
```